

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import re
```

```
In [2]: df=pd.read_csv('scaler_clustering.csv')
```

```
In [3]: df.head()
```

```
Out[3]:   Unnamed: 0  company_hash          email_hash  orgy
0           0  atrgxnnnt xzaxv  6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...  20'
              qtrxvzwt
1           1  xzegwgbb  b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...  20'
              rxbxnta
2           2  ojzwnvwnxw vx  4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...  20'
3           3  ngpgutaxv  effdede7a2e7c2af664c8a31d9346385016128d66bbc58...  20'
4           4  qxen sqghu  6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...  20'
```

```
In [4]: df.shape
```

```
Out[4]: (205843, 7)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 7 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        205843 non-null   int64  
 1   company_hash      205799 non-null   object  
 2   email_hash        205843 non-null   object  
 3   orgyear           205757 non-null   float64 
 4   ctc               205843 non-null   int64  
 5   job_position      153281 non-null   object  
 6   ctc_updated_year  205843 non-null   float64 
dtypes: float64(2), int64(2), object(3)
memory usage: 11.0+ MB
```

```
In [6]: df.describe()
```

Out[6]:

	Unnamed: 0	orgyear	ctc	ctc_updated_year
count	205843.000000	205757.000000	2.058430e+05	205843.000000
mean	103273.941786	2014.882750	2.271685e+06	2019.628231
std	59741.306484	63.571115	1.180091e+07	1.325104
min	0.000000	0.000000	2.000000e+00	2015.000000
25%	51518.500000	2013.000000	5.300000e+05	2019.000000
50%	103151.000000	2016.000000	9.500000e+05	2020.000000
75%	154992.500000	2018.000000	1.700000e+06	2021.000000
max	206922.000000	20165.000000	1.000150e+09	2021.000000

In [7]: df.describe(include=object)

Out[7]:

	company_hash	email_hash	job_positio
count	205799	205843	15328
unique	37299	153443	101
top	nvnv wgzohrnvwj otqcxwto	bbace3cc586400bbc65765bc6a16b77d8913836cf98b77c05488f02f5714a4b...	Backend Engineer
freq	8337	10	4355

In [8]: df=df[(df['orgyear'] < 2022) & (df['orgyear'] > 1900)]

Univariate Analysis

Non Visual Analysis

In [9]: df['email_hash'].value_counts()

Out[9]:

bbace3cc586400bbc65765bc6a16b77d8913836cf98b77c05488f02f5714a4b	10
3e5e49daa5527a6d5a33599b238bf9bf31e85b9efa9a94f1c88c5e15a6f31378	9
6842660273f70e9aa239026ba33bfe82275d6ab0d20124021b952b5bc3d07e6c	9
298528ce3160cc761e4dc37a07337ee2e0589df251d73645aae209b010210eee	9
d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf93246d4192a89d8065	8
..	
509c64729b4a325556a0e7e7af7fa12a1aa72405cf7367dd12019ddf88f036b8	1
fe49d1b2f7de7019e1773ae790669d9b5c1c7218a1849e748a87b19bc278e111	1
a1b8cb78f5fe704d7a4218f7e5438a12be71adc68a70ecfa12a38923f762f70	1
bf41860f81866a7d6f70f5d77f6dd1d35747918fa8ddaf6a4e564f37cd1cb337	1
0bcfc1d05f2e8dc4147743a1313aa70a119b41b30d4a1f7e738a6a87d3712c31	1

Name: email_hash, Length: 152301, dtype: int64

In [10]: df['company_hash'].value_counts()

```
Out[10]: nvnv wzohrvzj otqcxwto      8315
          xzegojo                  5362
          vbvkgz                   3461
          zgn vuurxwmrt vwwghzn    3275
          wgszxkvzn                 3230
          ...
          cvrthqyq                  1
          aqtvbch                  1
          xnxcxnx ucn rna          1
          nyt hzxctqoxnj ge ihttzorvza 1
          bvptbjnqxu td vbvkgz       1
          Name: company_hash, Length: 37073, dtype: int64
```

```
In [11]: df['job_position'].value_counts()
```


```
Out[11]: Backend Engineer           43416
          FullStack Engineer        24547
          Other                      17847
          Frontend Engineer         10351
          Engineering Leadership    6823
          ...
          Compliance auditor          1
          91                           1
          Senior Software Development Engineer (Backend) 1
          Messenger come driver       1
          Android Application developer 1
          Name: job_position, Length: 1015, dtype: int64
```

Visual Analysis

```
In [12]: sns.distplot(df['orgyear'])
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figur
e-level function with similar flexibility) or `histplot` (an axes-level f
unction for histograms).
    warnings.warn(msg, FutureWarning)
```


```
Out[12]: <AxesSubplot:xlabel='orgyear', ylabel='Density'>
```



```
In [13]: sns.distplot(df['ctc'])
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function).  
    warnings.warn(msg, FutureWarning)
```


```
Out[13]: <AxesSubplot:xlabel='ctc', ylabel='Density'>
```



```
In [14]: sns.distplot(df['ctc_updated_year'])
```


```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure-  
level function with similar flexibility) or `histplot` (an axes-level function).  
    warnings.warn(msg, FutureWarning)
```

```
Out[14]: <AxesSubplot:xlabel='ctc_updated_year', ylabel='Density'>
```




Bivariate Analysis

```
In [15]: df.hist(figsize=(10,10),bins=30);
```




```
In [16]: sns.pairplot(df, diag_kind='kde')
```

```
Out[16]: <seaborn.axisgrid.PairGrid at 0x7fc924510460>
```



```
In [17]: sns.heatmap(df.corr(), annot=True)
```


```
Out[17]: <AxesSubplot:>
```



Outlier Detection

```
In [18]: sns.boxplot(df['ctc'])
```

/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(
Out[18]: <AxesSubplot:xlabel='ctc'>




```
In [19]: sns.boxplot(df['orgyear'])
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
    warnings.warn(
```

```
Out[19]: <AxesSubplot:xlabel='orgyear'>
```




```
In [20]: sns.boxplot(df['ctc_updated_year'])
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
    warnings.warn(
```

```
Out[20]: <AxesSubplot:xlabel='ctc_updated_year'>
```



Feature Engineering and Missing value Imputation

```
In [21]: cat=df.describe(include=object).columns
```

```
In [22]: from sklearn.impute import SimpleImputer
si=SimpleImputer(strategy='most_frequent')
for i in cat:
    v=df[i].values
    df[i]=si.fit_transform(v.reshape(-1,1))
```

```
In [23]: from sklearn.impute import KNNImputer
ki=KNNImputer(n_neighbors=5)
c=df['orgyear'].values
df['orgyear']=ki.fit_transform(c.reshape(-1,1))
```

```
In [24]: df.dropna(axis=0, inplace=True)
```

```
In [25]: df.drop(['Unnamed: 0'], axis=1, inplace=True)
```

```
In [26]: df['job_position']=[re.sub('[^A-Za-z0-9 ]+', ' ', x) for x in df['job_posi
```

```
In [27]: df['company_hash']=[re.sub('[^A-Za-z0-9 ]+', ' ', x) for x in df['company_
```

```
In [28]: df['email_hash']=[re.sub('^[^A-Za-z0-9 ]+', ' ', x) for x in df['email_hash']]  
In [29]: df.drop_duplicates(inplace=True)  
In [30]: df['years_of_experience']=2023-df['orgyear']
```

Manual Clustering

Clustering using Company

```
In [31]: ch=df.groupby('company_hash')['ctc']  
In [32]: ch.mean()  
Out[32]: company_hash  
0 1000000.0  
0000 300000.0  
01 ojztqsj 550000.0  
05mz exzytvrny uqxcvnt rxbxnta 1100000.0  
1 100000.0  
...  
zyvzwt wgzohrnxz s tzsxztqo 940000.0  
zz 935000.0  
zzb ztdnstz vacxogqj ucn rna 600000.0  
zzgato 130000.0  
zzzbzb 720000.0  
Name: ctc, Length: 37073, dtype: float64  
In [33]: ch.median()  
Out[33]: company_hash  
0 100000.0  
0000 300000.0  
01 ojztqsj 550000.0  
05mz exzytvrny uqxcvnt rxbxnta 1100000.0  
1 100000.0  
...  
zyvzwt wgzohrnxz s tzsxztqo 940000.0  
zz 935000.0  
zzb ztdnstz vacxogqj ucn rna 600000.0  
zzgato 130000.0  
zzzbzb 720000.0  
Name: ctc, Length: 37073, dtype: float64  
In [34]: ch.max()
```

```
Out[34]: company_hash
0                      100000
0000                  300000
01 ojztqsj            830000
05mz exzytvrny uqxcvnt rxbxnta 1100000
1                      100000
...
zyvzwt wgzohrnxs tsxsztqo    940000
zz                     1370000
zzb ztdnstz vacxogqj ucn rna 600000
zzgato                 130000
zzzbzb                 720000
Name: ctc, Length: 37073, dtype: int64
```

```
In [35]: ch.min()
```

```
Out[35]: company_hash
0                      100000
0000                  300000
01 ojztqsj            270000
05mz exzytvrny uqxcvnt rxbxnta 1100000
1                      100000
...
zyvzwt wgzohrnxs tsxsztqo    940000
zz                     500000
zzb ztdnstz vacxogqj ucn rna 600000
zzgato                 130000
zzzbzb                 720000
Name: ctc, Length: 37073, dtype: int64
```

```
In [36]: ch.count()
```

```
Out[36]: company_hash
0                      2
0000                  1
01 ojztqsj            2
05mz exzytvrny uqxcvnt rxbxnta 1
1                      1
..
zyvzwt wgzohrnxs tsxsztqo    1
zz                     2
zzb ztdnstz vacxogqj ucn rna 2
zzgato                 1
zzzbzb                 1
Name: ctc, Length: 37073, dtype: int64
```

Clustering using Job Position

```
In [37]: jp=df.groupby('job_position')['ctc']
```

```
In [38]: jp.mean()
```

```
Out[38]: job_position
          650000.0
          SDE 2      1200000.0
          7          445000.0
    7033771951  100000000.0
    737         350000.0
          ...
student      1715000.0
support escalation engineer 2000000.0
system engineer      500000.0
system software engineer 610000.0
technology analyst     82000.0
Name: ctc, Length: 1003, dtype: float64
```

```
In [39]: jp.median()
```

```
Out[39]: job_position
          650000.0
          SDE 2      1200000.0
          7          445000.0
    7033771951  100000000.0
    737         350000.0
          ...
student      1715000.0
support escalation engineer 2000000.0
system engineer      500000.0
system software engineer 610000.0
technology analyst     82000.0
Name: ctc, Length: 1003, dtype: float64
```

```
In [40]: jp.max()
```

```
Out[40]: job_position
          700000
          SDE 2      1200000
          7          470000
    7033771951  100000000
    737         350000
          ...
student      2400000
support escalation engineer 2000000
system engineer      500000
system software engineer 610000
technology analyst     82000
Name: ctc, Length: 1003, dtype: int64
```

```
In [41]: jp.min()
```

```
Out[41]: job_position
          600000
          SDE 2      1200000
          7           420000
          7033771951 100000000
          737         350000
          ...
          student     1030000
          support escalation engineer 2000000
          system engineer      500000
          system software engineer 610000
          technology analyst    82000
Name: ctc, Length: 1003, dtype: int64
```

```
In [42]: jp.count()
```

```
Out[42]: job_position
          2
          SDE 2      1
          7           2
          7033771951 1
          737         1
          ..
          student     2
          support escalation engineer 1
          system engineer      1
          system software engineer 1
          technology analyst    1
Name: ctc, Length: 1003, dtype: int64
```

Clustering using Years of Experience

```
In [43]: yoe=df.groupby('years_of_experience')['ctc']
```

```
In [44]: yoe.mean()
```

```
Out[44]: years_of_experience  
2.0      4.174413e+06  
3.0      2.168744e+06  
4.0      1.661865e+06  
5.0      1.790543e+06  
6.0      2.014689e+06  
7.0      2.287597e+06  
8.0      2.412350e+06  
9.0      2.428622e+06  
10.0     2.360870e+06  
11.0     2.456325e+06  
12.0     2.177922e+06  
13.0     2.476629e+06  
14.0     2.633404e+06  
15.0     2.782173e+06  
16.0     2.754114e+06  
17.0     2.719412e+06  
18.0     2.848160e+06  
19.0     3.708461e+06  
20.0     3.641774e+06  
21.0     3.969004e+06  
22.0     6.701403e+06  
23.0     4.337689e+06  
24.0     3.842169e+06  
25.0     3.565805e+06  
26.0     2.937248e+06  
27.0     2.608812e+06  
28.0     3.577553e+06  
29.0     8.842375e+06  
30.0     5.440609e+06  
31.0     5.452064e+06  
32.0     7.771547e+06  
33.0     3.471514e+06  
34.0     2.520909e+06  
35.0     1.552200e+06  
36.0     2.496067e+06  
37.0     3.291250e+06  
38.0     2.924000e+06  
39.0     7.266667e+06  
41.0     9.400000e+05  
42.0     1.000000e+05  
44.0     3.100000e+06  
46.0     2.000000e+05  
47.0     8.000000e+05  
50.0     1.000000e+03  
51.0     2.300000e+06  
52.0     5.500000e+06  
53.0     9.700000e+05  
Name: ctc, dtype: float64
```

```
In [45]: yoe.median()
```

```
Out[45]: years_of_experience  
2.0      700000.0  
3.0      700000.0  
4.0      680000.0  
5.0      700000.0  
6.0      750000.0  
7.0      850000.0  
8.0      949999.0  
9.0      1000000.0  
10.0     1200000.0  
11.0     1360000.0  
12.0     1500000.0  
13.0     1600000.0  
14.0     1680000.0  
15.0     1750000.0  
16.0     2000000.0  
17.0     2100000.0  
18.0     2450000.0  
19.0     2500000.0  
20.0     2500000.0  
21.0     2700000.0  
22.0     2600000.0  
23.0     2750000.0  
24.0     3000000.0  
25.0     2900000.0  
26.0     2800000.0  
27.0     2600000.0  
28.0     2200000.0  
29.0     3000000.0  
30.0     2500000.0  
31.0     2300000.0  
32.0     2000000.0  
33.0     2700000.0  
34.0     1500000.0  
35.0     1300000.0  
36.0     2650000.0  
37.0     2750000.0  
38.0     2620000.0  
39.0     3000000.0  
41.0      950000.0  
42.0     100000.0  
44.0     3100000.0  
46.0     200000.0  
47.0     800000.0  
50.0      1000.0  
51.0     2300000.0  
52.0     5500000.0  
53.0     970000.0  
Name: ctc, dtype: float64
```

```
In [46]: yoe.max()
```

```
Out[46]: years_of_experience  
2.0      250000000  
3.0      200000000  
4.0      200000000  
5.0      255555555  
6.0      200000000  
7.0      200000000  
8.0      1000150000  
9.0      200000000  
10.0     200000000  
11.0     200000000  
12.0     200000000  
13.0     200000000  
14.0     110000000  
15.0     200000000  
16.0     200000000  
17.0     101500000  
18.0     100000000  
19.0     200000000  
20.0     190000000  
21.0     200000000  
22.0     100000000  
23.0     165000000  
24.0     100000000  
25.0     200000000  
26.0     16000000  
27.0     10500000  
28.0     100000000  
29.0     200000000  
30.0     100000000  
31.0     150000000  
32.0     200000000  
33.0     20000000  
34.0     19800000  
35.0     3600000  
36.0     5501400  
37.0     6500000  
38.0     4000000  
39.0     16300000  
41.0     1800000  
42.0     100000  
44.0     3100000  
46.0     200000  
47.0     800000  
50.0     1000  
51.0     2300000  
52.0     5500000  
53.0     1800000  
Name: ctc, dtype: int64
```

```
In [47]: yoe.min()
```

```
Out[47]: years_of_experience  
2.0          1000  
3.0           24  
4.0            16  
5.0           500  
6.0          1000  
7.0            15  
8.0          1000  
9.0             2  
10.0           6  
11.0          600  
12.0          1000  
13.0          1000  
14.0          1000  
15.0          2000  
16.0          1000  
17.0          1000  
18.0          1000  
19.0          1000  
20.0          1500  
21.0          1000  
22.0          1000  
23.0          2000  
24.0         13000  
25.0          1000  
26.0         30000  
27.0         21000  
28.0         25000  
29.0        135000  
30.0         24000  
31.0         62000  
32.0         65000  
33.0        90000  
34.0       100000  
35.0       180000  
36.0         40000  
37.0       500000  
38.0      2500000  
39.0      2500000  
41.0         60000  
42.0       100000  
44.0      3100000  
46.0       200000  
47.0       800000  
50.0          1000  
51.0      2300000  
52.0      5500000  
53.0       140000  
Name: ctc, dtype: int64
```

```
In [48]: yoe.count()
```

```
Out[48]: years_of_experience
2.0      3541
3.0     12781
4.0    22038
5.0    24068
6.0    22189
7.0    22061
8.0    19787
9.0    16115
10.0   11919
11.0   10131
12.0    7680
13.0   5550
14.0   3647
15.0   2653
16.0   2177
17.0   2015
18.0   1813
19.0   1423
20.0    996
21.0    668
22.0    701
23.0    491
24.0    338
25.0    279
26.0    232
27.0    133
28.0     94
29.0     64
30.0     69
31.0     47
32.0     75
33.0     37
34.0     22
35.0     10
36.0      6
37.0      8
38.0      5
39.0      3
41.0      4
42.0      1
44.0      1
46.0      1
47.0      1
50.0      1
51.0      1
52.0      1
53.0      2
Name: ctc, dtype: int64
```

```
In [49]: df1=df.groupby(['company_hash','years_of_experience'])['ctc'].mean()
```

Merging the same with original dataset carefully and creating some flags showing learners with CTC greater than the Average of their Company's department having same Years of Experience - Call that flag designation with values [1,2,3]

```
In [50]: df2=df.merge(df1, left_on=['company_hash', 'years_of_experience'], right_on=[])
In [51]: df2['ctc_avg']=df2['ctc'].round(2)
In [52]: df2['designation']=np.where(df2['ctc']>df2['ctc_avg'],1,10)
         df2['designation']=np.where(df2['ctc']<df2['ctc_avg'],3,df2['designation'])
         df2['designation']=np.where(df2['ctc']==df2['ctc_avg'],2,df2['designation'])
In [53]: df2.drop('ctc_avg', axis=1, inplace=True)
```

Doing above analysis at Company & Job Position level. Name that flag Class with values [1,2,3]

```
In [54]: df1=df.groupby(['company_hash', 'job_position'])['ctc'].mean()
In [55]: df2=df2.merge(df1, left_on=['company_hash', 'job_position'], right_on=['comp
In [56]: df2['ctc_avg']=df2['ctc'].round(2)
In [57]: df2['class']=np.where(df2['ctc']>df2['ctc_avg'],1,10)
         df2['class']=np.where(df2['ctc']<df2['ctc_avg'],3,df2['class'])
         df2['class']=np.where(df2['ctc']==df2['ctc_avg'],2,df2['class'])
In [58]: df2.drop('ctc_avg', axis=1, inplace=True)
```

Repeating the same analysis at the Company level. Name that flag Tier with values [1,2,3]

```
In [59]: df1=df.groupby(['company_hash'])['ctc'].mean()
In [60]: df2=df2.merge(df1, left_on=['company_hash'], right_on=['company_hash'], suff
In [61]: df2['ctc_avg']=df2['ctc'].round(2)
In [62]: df2['tier']=np.where(df2['ctc']>df2['ctc_avg'],1,10)
         df2['tier']=np.where(df2['ctc']<df2['ctc_avg'],3,df2['tier'])
         df2['tier']=np.where(df2['ctc']==df2['ctc_avg'],2,df2['tier'])
```

Based on the manual clustering done so far, answering few questions

Top 10 employees (earning more than most of the employees in the company) - Tier 1

```
In [63]: df2[df2['tier']==1].sort_values(by=['ctc'], ascending=False).head(10)
```

	company_hash	email_hash	orgyear
150927	obvqnuqxdwgb	5b4bed51797140db4ed52018a979db1e34cee49e27b488...	2018.0
7751	gnytqo	e7b2b159325f77e7abd5aa938371d7f9425530b36e703f...	2013.0
16317	ntwy bvyxzaqv	26946b7b12b7daa5b4f025d7c4c5c3ee0fbdb3bbe01356...	2016.0
80868	ovrtoegqwt	eb552f9d6f12d47656472a3f7c6a6625ebf3d699edb4b0...	2013.0
27284	vbkpgz	9e785d33821db67c01becc1c36f901d79d3142c1d13bd8...	2017.0
16581	ntwy bvyxzaqv	7a723f5b71698674b79bd2195c3bb58d3fcf4ddb75a04...	2019.0
106973	zvz	0f7e7ebdae89364a5c20d32fed1886003a4375ed25531c...	2014.0
62012	ovu	a35a5abbe9fb056421bdd9aca4440acfb93e37c823564d...	2017.0
67124	ntrtutqeqqbvzwt	979d02840c45c1d5790306130a0977aab05f2bd2679687...	2013.0
145249	vznxzg rvmo	634fd283565b8954513a6ad0e47cedb0fa8847923149fb...	2019.0

Top 10 employees of data science in Amazon / TCS etc earning more than their peers - Class 1

```
In [64]: df2[(df2['class']==1)& (df2['job_position']=='Data Scientist')|(df2['job_
```

Out [64]:

	company_hash	email_hash	orgyear
148305	ihvaqvnwx xzoxsyno ucn rna	bd222ea783ee372da4e0ad60fdccce0b8f37999a032025...	2015.0
72609	mqxonrtwgzt v bvyxzaqv sqghu wgbuvzj	cda8d723438e81185d2ee8c348870a4612eea974cdb2db...	2017.0
16569	ntwy bvyxzaqv	6ad86d120e39db485331f9a0b2b1f15ce2a7bdaee778ab...	2021.0
144225	xzzgcv ogrhnxgzo	6b6dd66bae787dd4dd417e1777f8ea5a057257e9019995...	2016.0
126143	ptzgbt	4ddef8762b7585c6ee7b8c06834778f3aa00eb3be312b0...	2020.0
151738	wgqt wgbutnt	75f5b46d47310c3923e93329a62a1aa78d478803f0a685...	2016.0
14369	zgn vuurxwvrmr vwvghzn	544e75b477f8644eb71281133c62c19732547837e80e51...	2021.0
69624	zvsvqqg	15adaeb2eef9c0ee8a0f18e189bf426be390f5d1e911fd...	2021.0
2267	ztfstz ogenfvqt	3c64901d83458f3b7b8eed6fb529ee3a4c14d49339c398...	2017.0
52722	bgqsvz onvzrtj	2bede29959707d8c6f283d98319361c386baa6fa5c8028...	2021.0

Bottom 10 employees of data science in Amazon / TCS etc earning less than their peers - Class 3

In [65]: df2[(df2['class']==3)& (df2['job_position']=='Data Scientist')|(df2['job_

Out [65]:

		company_hash	email_hash	orgyear
15788	ytfrtnn uvwpvqa tzntquqxot	8274b3188470cd1c4914e7face490111e27f239457e62d...		2018.0
104395	sggsrt	fb64af615420e06d46a1965f59068b34460fb3cbe70541...		2018.0
34709	uvjovet sqghu	3cc0c85d198d0e56a4cdefb6496333f59b97f87c293262...		2018.0
138449	vqxosrgmvr	3675f79c7e05de96ccf189c818b84b487cb1aa3f6b80e8...		2015.0
44901	nvnv wgzohrnvzwj otqcxwto	3175d03fd4618eb293d6f5a1d13d42a0c79f68e9acaaa3...		2020.0
145930	exznqhon ogrhnxgzo ucn rna	ab2dc9db23c3104f0b6b3dbd4cdd5bfb9e5829b8b7943d...		2017.0
143305	ovbohzs trtwnqg btwyvxwo	e374eea75640881206a21894f69190138c2c0535277dc1...		2017.0
90145	onhatzn	bd9c04a574090e05b366a81cdb2f3f565d0c60fa8b1647...		2021.0
108835	bxyhu wgbbhzxwvnxz	690f6fdab1ab7514a6a9325ebd6cfe910dbf12d46b6fde...		2018.0
35910	srgmvrast xznrrxstzwt ge nyxzso	8001bc017fbe95541d23f5780c3edb988b7d9b2225e39e...		2017.0

Bottom 10 employees (earning less than most of the employees in the company)- Tier 3

In [66]:

```
df2[df2['tier']==3].sort_values(by=['ctc'], ascending=False).tail(10)
```

Out [66]:

		company_hash	email_hash	orgyear
12503		zgn vuurxwvmlrt vwvghzn	8d1e069a03fc437876b406b8c93bc7e07577f9836222bd...	2021.0
116155	mtznrtj ojontbo		7c8e0d8194db4deb41cbc9b3b6c428e0f9ab289436638e...	2016.0
45710		nvvv wgzohrnvwj otqcxwto	8625d6d072e12dad0c5748ab010e1d0315736a359e2bb5...	2013.0
10048	xb v onhatzn		4eea97c023bd58395edce18538831df9a735180f88f79d...	2020.0
44446		nvvv wgzohrnvwj otqcxwto	80ba0259f9f59034c4927cf3bd38dc9ce2eb60ff18135b...	2012.0
71341		gjg	b995d7a2ae5c6f8497762ce04dc5c04ad6ec734d70802a...	2018.0
192736		xm	b8a0bb340583936b5a7923947e9aec21add5ebc50cd60b...	2016.0
101949	xzntqcxtfmxn		23ad96d6b6f1ecf554a52f6e9b61677c7d73d8a409a143...	2013.0
101950	xzntqcxtfmxn		f2b58aeed3c074652de2cf3c0717a5d21d6fbef342a78...	2013.0
101946	xzntqcxtfmxn		3505b02549ebe2c95840ac6f0a35561a3b4cbe4b79cdb1...	2014.0

Top 10 employees in Amazon- X department - having 5/6/7 years of experience earning more than their peers - Tier X

In [67]: df2[(df2['designation']==1) & (df2['years_of_experience']==5) | (df2['years_

Out [67]:

	company_hash		email_hash	orgyear
150927	obvqnuqxdwgb	5b4bed51797140db4ed52018a979db1e34cee49e27b488...		2018.0
7787	gnytqo	06c7cd6c1a50803bac81950ad450a0e78165cb7c09ddfe...		2018.0
40846	nvnv wgzohrnvwj otqcxwto	10c67fa43291396c3f72c9ac34b99a6d9fb2c2007e6964...		2018.0
38310	nvnv wgzohrnvwj otqcxwto	979235a69267e855c0361f670e5941138307caf43fa986...		2018.0
26901	vbvkgz	8dfe6251bd4ec533f02ddceb98b3dcebb9550ccd4ef2e6...		2018.0
38312	nvnv wgzohrnvwj otqcxwto	97d25613e7bc3f47c87492d311f77232c105e4bc9ce642...		2018.0
17404	xzegojo	d418a571a2cf09cb8cbd9eceade19381777835f58d8f98...		2018.0
45892	nvnv wgzohrnvwj otqcxwto	1b95e7ba0ee82100ca5a034239fa0203a1bec14280b82a...		2018.0
56196	xwxwx mvzp	2311bf023218afe93d650cac03abb7a40f7fa55c08d260...		2018.0
128470	nqvctrnqxvzsrt	3eb73d5c74bfd5f7a195d7f56e4922a46357c6ae0a8732...		2018.0

Top 10 companies (based on their CTC)

In [68]: `a=df2.groupby('company_hash')['ctc_avg'].mean()`

In [69]: `a=a.sort_values(ascending=False).head(10)`

In [70]: `a`

Out [70]:

company_hash	ctc_avg
whmxw rgsxwo uqxcvnt rxbxnta	1.000150e+09
aveegaxr xzntqzvnxzvr hzxctqoxnj	2.500000e+08
uhxoovzwt xn vacxogqj vza exzvzxvr otqcxwto rru	2.000000e+08
wvquvzntq	2.000000e+08
durgfxk ogrhnxgzo	2.000000e+08
qn	2.000000e+08
ztbyvzo ogrhnxgzo ucn rna	2.000000e+08
ihvrxnvo srgmvr rru	2.000000e+08
bvtonqg wxcxr vza xzntqzgqo	2.000000e+08
egd z wrgha ntwyzrgsxt	2.000000e+08
Name: ctc_avg, dtype: float64	

Top 2 positions in every company (based on their CTC)


```
In [71]: df_grouped = df.groupby(['job_position']).apply(lambda x: x.sort_values('ctc'))  
  
In [72]: df_grouped.groupby('company_hash')['ctc'].nlargest(2)  
  
Out[72]:

| company_hash                 | job_position                           |        |        |
|------------------------------|----------------------------------------|--------|--------|
| 0                            | Backend Engineer                       | 2940   | 100000 |
|                              | Other                                  | 16824  | 100000 |
| 0000                         | Other                                  | 197540 | 300000 |
| 01_ojztqsj                   | Frontend Engineer                      | 55241  | 830000 |
|                              | Android Engineer                       | 74429  | 270000 |
|                              |                                        |        | ...    |
| zz                           | Backend Engineer                       | 14670  | 500000 |
| zzb_ztdnstz_vacxogqj_ucn_rna | Backend Engineer                       | 146629 | 600000 |
|                              | FullStack Engineer                     | 72983  | 600000 |
| zzgato                       | Backend Engineer                       | 117023 | 130000 |
| zzzbzb                       | Other                                  | 15838  | 720000 |
|                              |                                        |        |        |
|                              | Name: ctc, Length: 50181, dtype: int64 |        |        |


```

Data processing for Unsupervised clustering - Label encoding/ One- hot encoding, Standardization of data

```
In [73]: from category_encoders import CountEncoder  
ce=CountEncoder()  
df['email_hash']=ce.fit_transform(df['email_hash'])  
  
In [74]: from category_encoders import CountEncoder  
ce=CountEncoder()  
df['company_hash']=ce.fit_transform(df['company_hash'])  
  
In [75]: from category_encoders import CountEncoder  
ce=CountEncoder()  
df['job_position']=ce.fit_transform(df['job_position'])  
  
In [76]: sns.pairplot(df)  
  
Out[76]: <seaborn.axisgrid.PairGrid at 0x7fc8d0c3a9d0>
```



```
In [77]: num_cols=df._get_numeric_data().columns
```

```
In [78]: '''def outliers(data,feature):
    q1=data[feature].quantile(0.05)
    q3=data[feature].quantile(0.95)
    iqr=q3-q1
    ul=q3+1.5*iqr
    ll=q1-1.5*iqr
    return ul,ll'''
```

```
Out[78]: 'def outliers(data,feature):\n    q1=data[feature].quantile(0.05)\n    q3=data[feature].quantile(0.95)\n    iqr=q3-q1\n    ul=q3+1.5*iqr\n    ll=q1-1.5*iqr\n    return ul,ll'
```

```
In [79]: '''for i in num_cols:
    ul,ll=outliers(df,i)
    df=df[(df[i]<ul) & (df[i]>ll)]'''
```

```
Out[79]: 'for i in num_cols:\n    ul,ll=outliers(df,i)\n    df=df[(df[i]<ul) & (df[i]>ll)]'
```

Outlier Removal using Local Outlier Factor

```
In [80]: from sklearn.neighbors import LocalOutlierFactor
clf = LocalOutlierFactor(n_neighbors=25, contamination=0.05)
is_o=clf.fit_predict(df)
```

```
In [81]: df['lof']=is_o
df
```

Out[81]:

	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year
0	9	2	2016.0	1100000	17847	2020.0
1	401	2	2018.0	449999	24541	2019.0
2	1	2	2015.0	2000000	86961	2020.0
3	68	1	2017.0	700000	86961	2019.0
4	6	2	2017.0	1400000	24541	2019.0
...
205838	17	2	2008.0	220000	86961	2019.0
205839	105	1	2017.0	500000	86961	2020.0
205840	159	1	2021.0	700000	86961	2021.0
205841	968	1	2019.0	5100000	86961	2019.0
205842	533	1	2014.0	1240000	86961	2016.0

195879 rows × 8 columns

```
In [82]: df=df[df['lof']==1]
```

```
In [83]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X=sc.fit_transform(df)
```

Checking clustering tendency

```
In [84]: from pyclustertend import hopkins
hopkins(X, X.shape[0])
```


Out[84]: 0.003810942596027236

Elbow Method

```
In [85]: num_clusters=np.arange(1,10)
```

```
In [86]: from sklearn.cluster import KMeans
inertia=[]
for i in num_clusters:
    km=KMeans(n_clusters=i)
    km.fit(X)
    inertia.append(km.inertia_)
```

```
In [87]: plt.plot(num_clusters,inertia)
plt.scatter(num_clusters,inertia)
plt.grid()
plt.show()
```



K-means clustering

```
In [88]: km=KMeans(n_clusters=2)
km.fit(X)
```

```
Out[88]: KMeans(n_clusters=2)
```

```
In [89]: pred=km.fit_predict(X)
```

```
In [90]: clusters = pd.DataFrame(X, columns=df.columns)
clusters['label'] = km.labels_
clusters.head(3)
```


	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year	ye
0	-0.484794	0.561983	0.224445	-0.099497	-0.710530	0.303744	
1	-0.264028	0.561983	0.698980	-0.157010	-0.534433	-0.451547	
2	-0.489299	0.561983	-0.012822	-0.019863	1.107632	0.303744	

```
In [91]: np.unique(km.labels_)
```

```
Out[91]: array([0, 1], dtype=int32)
```

```
In [92]: import plotly.express as px
```

```
fig = px.scatter_3d(clusters, x='company_hash', y='years_of_experience',
fig.update_traces(marker=dict(size=2,
                                line=dict(width=2,
                                          color='DarkSlateGrey')),
                                selector=dict(mode='markers')))
fig.show()
```



Hierarchical clustering

```
In [93]: q=pd.DataFrame(X)
sam=q.sample(20000).values
```


```
In [94]: from sklearn.cluster import AgglomerativeClustering
hc = AgglomerativeClustering(n_clusters=3, affinity = 'euclidean', linkage='ward')
y_pred = hc.fit_predict(sam)
```

```
In [95]: clusters = pd.DataFrame(sam, columns=df.columns)
clusters['label'] = hc.labels_
clusters.head(3)
```


```
Out[95]:   company_hash  email_hash  orgyear      ctc  job_position  ctc_updated_year  ye
0      -0.380042    0.561983 -0.012822 -0.099497     1.107632        1.059034
1      1.366938   -0.684951  0.936248 -0.072952    -0.534433        0.303744
2      -0.484794   -0.684951 -0.487357  0.006682     1.107632       -3.472708
```

```
In [96]: def viz_clusters(hc):
    plt.scatter(clusters['ctc'], clusters['job_position'], c=clusters['label'])
    plt.xlabel('amount_spent')
    plt.ylabel('amount_discount')
    plt.show()

viz_clusters(hc)
```



```
In [97]: import plotly.express as px  
  
fig = px.scatter_3d(clusters, x='company_hash', y='ctc', z='job_position'  
fig.update_traces(marker=dict(size=2), selector=dict(mode='markers'))  
fig.show()
```



In []:

In []: